

COL 718 - LAB2

# Simulating a system with 2 caches(L1 and L2) for Different Parameters

Submitted to: Prof. Kolin Paul  
IIT Delhi

Madhur Aswani  
2023JVL2241

August 14, 2023

## 1 Introduction

In this LAB I simulated a system with CPU Type *Timing* and different Cache parameters operating at a fixed frequency of *600MHz*. The aim of this lab is to see how different parameters of the cache affect the system.

## 2 Experimental Details

In this Experiment, I have added two L1 Caches and one L2 Cache. The 2 L1 Caches are **L1ICache** and **L1DCache** each specifically connected to CPU *ICache* port and *DCache* port respectively. The CPU sends the Instruction requests to **L1ICache** and data requests **L1DCache**. For this experiment, the size of **L1ICache** has been fixed at *16kB*. In comparison, the other cache parameters are varied equally over both the L1Caches.

### 2.1 Parameters of Experiment

We are going to vary the below parameters in this experiment -

- **L1nL2ratio** - It refers to the ratio of sizes of L1 and L2 cache in the system [2, 4, 8, 16].
- **L1Latency** - It is the Latency of L1 Cache in num of cycles [1-5].
- **L2Latency** - It is the Latency of L2 Cache in num of cycles [10, 15, 20, 25, 30].
- **L1assoc** - It is the associativity of L1 Cache [2, 4, 8].
- **L2assoc** - It is the associativity of L2 Cache [4, 8, 16].
- **L1Size** - It is the Size of L1 Cache ['64kB', '128kB'].
- **blkSize** - When given 0 it uses unblocked MM and is varied for 4 and 8 block sized MM.

## 2.2 Stats Focused

In this experiment, I will be focusing on the below statistics -

- **simSeconds** - This stat gives the simulated seconds of the system.
- **hostInstRate** - This state gives the simulator instruction rate (inst/s) ((Count/Second))
- **numCycles** - Number of CPU cycles simulated (Cycle)
- **cpu.cpi** - CPI: cycles per instruction (core level) ((Cycle/Count))
- **cpu.ipc** - IPC: instructions per cycle (core level) ((Count/Cycle))
- **dcache.overallMissRate::total** - The stat gives the Overall Miss Rate of *L1DCache*.
- **dcache.overallAvgMissLatency::total** - The stat gives the Overall Avg Miss Latency of *L1DCache*.
- **icache.overallMissRate::total** - The stat gives the Overall Miss Rate of *L1ICache*.
- **icache.overallAvgMissLatency::total** - The stat gives the Overall Avg Miss Latency of *L1Cache*.
- **l2cache.overallMissRate::total** - The stat gives the Overall Miss Rate of *L2Cache*.
- **l2cache.overallAvgMissLatency::total** - The stat gives the Overall Avg Miss Latency of *L2Cache*.

## 2.3 Methodology

In this experiment, I have created a system with a *X86TimingSimpleCPU* CPU and 2 L1 Catches, one for Instruction and one for Data. Apart from this, there is an L2 Cache. Everything is specified in the *Config.py* file in Appendix A. There is also a *Cache.py* which defines the cache settings. For workload, I am using MM of  $25 \times 25$  Matrix and have added blocksize as an argument to the binary generated. The code is given in Appendix A. I have also written a *bash.sh* file to automate the simulation for different parameters and store the generated stat files.

## 3 Results

After running the bash file, I obtained about 288 stat files and created a script that writes the stats listed in these stat files to a combined *Data.csv* file. I have obtained some interesting plots from the file and observed the below.

- As the *blkSize* Increases the *dcache.overallMissRate::total* decreases with the unblocked MM having the highest missing rate. There is a significant decrease when we start using blocked MM.
- The plot of *hostInstRate* with *blockSize* for different associativity of l1, l2 and size of l1 cache show changed behavior. For an *l1assoc* of 2, *l2assoc* of 8 and *l1Size* of 64kB, the unblocked MM performs the best with the highest *hostInstRate*. As we start using Blocked MM, it still remains the highest for a block size of 4 but on using eight block size the configuration of *l1assoc* of 4 and *l2assoc* of 16 has the highest *hostInstRate*. Whereas for 128KB of cache size the *l1assoc* of 4 and *l2assoc* of 16 had the highest *hostInstRate*, but if






Figure 1: A Lineplot of blkSize and dcacheMissRate



(a) A Lineplot of blkSize and hostInstRate for 64kB (b) A Lineplot of blkSize and hostInstRate for 128kB

we move to block size of 8, we observe that the configuration of *l1assoc* of 4 and *l2assoc* of 8 has the highest *hostInstRate*.

- From the above, it seems that the *simSeconds* should decrease as the block size increases, but if you plot them, the entire opposite is observed. I guess the number of operations is high in case of blocked due to 6 nested for loops in the code.




Figure 3: A Lineplot of blkSize and simSeconds

## A Code Files

### Config.py

---

```
import m5
from m5.objects import *
import argparse
from Cache import *

parser = argparse.ArgumentParser(description='Config For Simulation')
parser.add_argument("-f", "--freq", type=str)
parser.add_argument("--l1Size", type=str)
parser.add_argument("--l2Ratio", type=float)
parser.add_argument("--l1assoc", type=int)
parser.add_argument("--l2assoc", type=int)
parser.add_argument("--blksize", type=str)
parser.add_argument("--l1Latency", type=int)
parser.add_argument("--l2Latency", type=int)

args=parser.parse_args()
size_l1=args.l1Size
l1s=int(size_l1[:-2])
l2s=l1s*args.l2Ratio
l2Size=str(int(l2s))+ "kB"
system = System()
l1Icache=L1Cache(args.l1assoc,args.l1Latency,size='16kB')
l1Dcache=L1Cache(args.l1assoc,args.l1Latency,size=args.l1Size)
l2Cache=L2Cache(assoc=args.l2assoc,latency=args.l2Latency,size=l2Size)

system.clk_domain = SrcClockDomain()
system.clk_domain.clock = args.freq
system.clk_domain.voltage_domain = VoltageDomain()

system.mem_mode = 'timing'
system.mem_ranges = [AddrRange('1GB')]

system.cpu = X86TimingSimpleCPU()

system.cpu.icache = l1Icache
system.cpu.dcache = l1Dcache

system.cpu.icache.cpu_side=system.cpu.icache_port
system.cpu.dcache.cpu_side=system.cpu.dcache_port

system.l2bus = L2XBar()

system.cpu.icache.connectBus(system.l2bus)
```

```

system.cpu.dcache.connectBus(system.l2bus)

system.l2cache = L2Cache
system.l2cache.connectCPUSideBus(system.l2bus)
system.membus = SystemXBar()
system.l2cache.connectMemSideBus(system.membus)

system.cpu.createInterruptController()
system.cpu.interrupts[0].pio = system.membus.mem_side_ports
system.cpu.interrupts[0].int_requestor = system.membus.cpu_side_ports
system.cpu.interrupts[0].int_responder = system.membus.mem_side_ports

system.system_port = system.membus.cpu_side_ports

system.mem_ctrl = MemCtrl()
system.mem_ctrl.dram = DDR3_1600_8x8()
system.mem_ctrl.dram.range = system.mem_ranges[0]
system.mem_ctrl.port = system.membus.mem_side_ports

system.workload = SEWorkload.init_compatible("../Github/IITD_Learning/Architecture
of HPC/LAB2/loadBinary")

process = Process()
process.cmd=["../Github/IITD_Learning/Architecture of
HPC/LAB2/loadBinary",args.blksize]
system.cpu.workload = process
system.cpu.createThreads()
root = Root(full_system = False, system = system)
m5.instantiate()

print("Beginning simulation!")
exit_event = m5.simulate()

print('Exiting @ tick {} because {}'.format(m5.curTick(), exit_event.getCause()))

```

---

## Cache.py

---

```

from m5.objects import Cache

class L1Cache(Cache):
    assoc = 2
    data_latency = 2
    tag_latency = 2
    response_latency = 2
    mshrs = 4
    tgts_per_mshr = 20
    size='64kB'
    def __init__(self,assoc,latency,size):
        super(L1Cache, self).__init__()
        self.assoc=assoc
        self.data_latency=latency
        self.tag_latency=latency
        self.response_latency=latency

```

```

        self.size=size
    def connectBus(self, bus):
        self.mem_side = bus.cpu_side_ports

class L2Cache(Cache):
    assoc = 8
    tag_latency = 20
    data_latency = 20
    response_latency = 20
    mshrs = 20
    tgts_per_mshr = 12

    size='256kB'
    def __init__(self,assoc,latency,size):
        super(L2Cache, self).__init__()
        self.assoc=assoc
        self.data_latency=latency
        self.tag_latency=latency
        self.response_latency=latency
        self.size=size
    def connectCPUSideBus(self, bus):
        self.cpu_side = bus.mem_side_ports

    def connectMemSideBus(self, bus):
        self.mem_side = bus.cpu_side_ports

```

---

## load.c

---

```

#include <stdio.h>

int main(int argc, char* argv[]){

    int res[25][25];
    int sum;
    int block=atoi(argv[1]);

    int a[25][25]={1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20,
        21, 22, 23, 24, 25}, {1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17,
        18, 19, 20, 21, 22, 23, 24, 25}, {1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14,
        15, 16, 17, 18, 19, 20, 21, 22, 23, 24, 25}, {1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11,
        12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23, 24, 25}, {1, 2, 3, 4, 5, 6, 7, 8,
        9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23, 24, 25}, {1, 2, 3, 4, 5,
        6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23, 24, 25}, {1, 2,
        3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23, 24,
        25}, {1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21,
        22, 23, 24, 25}, {1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18,
        19, 20, 21, 22, 23, 24, 25}, {1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15,
        16, 17, 18, 19, 20, 21, 22, 23, 24, 25}, {1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12,
        13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23, 24, 25}, {1, 2, 3, 4, 5, 6, 7, 8, 9,
        10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23, 24, 25}, {1, 2, 3, 4, 5, 6,

```





```

    }
    res[i][j]=sum;
}

}
}
else{
    printf("Doing Blocked");
    int bi=0;
    int bj=0;
    int bk=0;
    int i=0;
    int j=0;
    int k=0;
    int blockSize=block;
    int n=25;

    for(bi=0; bi<n; bi+=blockSize)
        for(bj=0; bj<n; bj+=blockSize)
            for(bk=0; bk<n; bk+=blockSize)
                for(i=0; i<blockSize; i++)
                    for(j=0; j<blockSize; j++)
                        for(k=0; k<blockSize; k++){
                            res[bi+i][bj+j] += a[bi+i][bk+k]*b[bk+k][bj+j];
                        }
                }
        }

printf("HelloWorld");
}

```

---

## bash.sh

---

```

#!/bin/bash
size=("64kB" "128kB")
ratio=(2 4 8 16)
l1assoc=(2 4 8)
l2assoc=(4 8 16)
for n in ${size[@]};
do
    for r in ${ratio[@]}
    do
        for a in ${l1assoc[@]}
        do
            for b in ${l2assoc[@]}
            do
                for l1 in {1..5}
                do
                    for l2 in {10..30..5}

```

```

do
./build/X86/gem5.opt ../../Github/IITD_Learning/Architecture\ of\
HPC/LAB2/Config.py -f "600MHz" --l1Size $n --l2Ratio $r
--l1assoc $a --l2assoc $b --blksize 1 --l1Latency $l1
--l2Latency $l2
cp m5out/stats.txt ../../Github/IITD_Learning/Architecture\ of\
HPC/LAB2/stats/stats_l1Size_$n-l1n12ratio_$r-l1assoc_$a-l2assoc_$b-l1Latency_$l1-l2Latency_$l2
done
done
done
done
done

```

---